

Base Template which will be saving the data

```
# =====
# Deep Learning Programming Assignment 2 (Keras / Colab)
# Part 1: CIFAR-10 (keras.datasets)
# Part 2: CIFAR-100 (keras.datasets)
# Part 3: Animal10 (image_dataset_from_directory, 20% split)
# Saves: best/last full model (.keras) + best/last weights (.weights.h5) + logs
# =====

# -----
# 0) CHOOSE PART + RUN NAME
# -----
PART = 1 # <-- set to 1, 2, or 3
RUN_NAME = f"part{PART}_run1" # change for new experiments (run2, run3...)

# For Part 3 dataset path (your folder structure)
ANIMAL_DIR = "/content/drive/MyDrive/AnimalsData/Animals10/raw-img"
```

1. MOUNT DRIVE + FOLDERS

```
# -----
# 1) Mount Drive + folders
# -----
from google.colab import drive
drive.mount('/content/drive')

import os, json, random
import numpy as np
import pandas as pd
import tensorflow as tf
import matplotlib.pyplot as plt

PROJECT_DIR = "/content/drive/MyDrive/DL_Prog2"
RUN_DIR = os.path.join(PROJECT_DIR, RUN_NAME)

CKPT_DIR = os.path.join(RUN_DIR, "checkpoints")
LOG_DIR = os.path.join(RUN_DIR, "logs")
OUT_DIR = os.path.join(RUN_DIR, "outputs")

for d in [CKPT_DIR, LOG_DIR, OUT_DIR]:
    os.makedirs(d, exist_ok=True)

print("RUN_DIR:", RUN_DIR)
print("TF:", tf.__version__)
print("GPU:", tf.config.list_physical_devices('GPU'))

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True)
RUN_DIR: /content/drive/MyDrive/DL_Prog2/part1_run1
TF: 2.19.0
GPU: [PhysicalDevice(name='/physical_device:GPU:0', device_type='GPU')]
```

2. CONFIG + REPRODUCIBILITY

```
# -----
# 2) Config + reproducibility
# -----
SEED = 42
random.seed(SEED)
np.random.seed(SEED)
tf.random.set_seed(SEED)

# Recommended defaults (edit if needed)
CONFIG = {
    "part": PART,
    "seed": SEED,
```

```

    "epochs": 30,
    "batch_size": 64,
    "lr": 1e-3,
    "val_split": 0.2,      # Part 3 requires 20% testing/validation
    "img_size_part3": 96,  # 64 faster; 128 potentially better
}

with open(os.path.join(RUN_DIR, "config.json"), "w") as f:
    json.dump(CONFIG, f, indent=2)

```

CONFIG

```

{'part': 1,
 'seed': 42,
 'epochs': 30,
 'batch_size': 64,
 'lr': 0.001,
 'val_split': 0.2,
 'img_size_part3': 96}

```

3. HELPER FUNCTIONS

```

# -----
# 3) Helper: show samples + plotting
# -----
def show_samples_np(x, y, n=12):
    idx = np.random.choice(len(x), n, replace=False)
    plt.figure(figsize=(12,4))
    for i, k in enumerate(idx):
        plt.subplot(2, n//2, i+1)
        plt.imshow(x[k])
        plt.title(str(int(y[k])))
        plt.axis("off")
    plt.tight_layout()
    plt.show()

def plot_accuracy_from_csv(csv_path, title="Accuracy vs Epoch"):
    df = pd.read_csv(csv_path)
    epochs_axis = np.arange(len(df))
    plt.figure()
    if "accuracy" in df.columns:
        plt.plot(epochs_axis, df["accuracy"], label="train_acc")
    if "val_accuracy" in df.columns:
        plt.plot(epochs_axis, df["val_accuracy"], label="val_acc")
    plt.xlabel("Epoch")
    plt.ylabel("Accuracy")
    plt.title(title)
    plt.legend()
    plt.grid(True)
    plt.show()

```

4. DATA AUGMENTATION + MODEL BUILDERS

```

# -----
# 4) Data augmentation + model builders
# -----
data_augmentation = tf.keras.Sequential([
    tf.keras.layers.RandomFlip("horizontal"),
    tf.keras.layers.RandomRotation(0.05),
    tf.keras.layers.RandomZoom(0.1),
    tf.keras.layers.RandomContrast(0.1),
])

def build_cnn(input_shape, num_classes, lr):
    inputs = tf.keras.Input(shape=input_shape)

    x = data_augmentation(inputs)
    x = tf.keras.layers.Rescaling(1./255)(x) # expects uint8 images

```

```

x = tf.keras.layers.Conv2D(32, 3, padding="same", activation="relu")(x)
x = tf.keras.layers.MaxPool2D()(x)

x = tf.keras.layers.Conv2D(64, 3, padding="same", activation="relu")(x)
x = tf.keras.layers.MaxPool2D()(x)

x = tf.keras.layers.Conv2D(128, 3, padding="same", activation="relu")(x)
x = tf.keras.layers.MaxPool2D()(x)

x = tf.keras.layers.Dropout(0.3)(x)
x = tf.keras.layers.Flatten()(x)
x = tf.keras.layers.Dense(256, activation="relu")(x)
x = tf.keras.layers.Dropout(0.4)(x)

outputs = tf.keras.layers.Dense(num_classes, activation="softmax")(x)

model = tf.keras.Model(inputs, outputs)
model.compile(
    optimizer=tf.keras.optimizers.Adam(learning_rate=lr),
    loss=tf.keras.losses.SparseCategoricalCrossentropy(),
    metrics=["accuracy"]
)
return model

```

5. SAVING + RESUME LOGIC

```

# -----
# 5) Saving + resume logic (works for all parts)
# -----
best_full = os.path.join(CKPT_DIR, "best.keras")
last_full = os.path.join(CKPT_DIR, "last.keras")
best_wts = os.path.join(CKPT_DIR, "best.weights.h5")
last_wts = os.path.join(CKPT_DIR, "last.weights.h5")

csv_log = os.path.join(LOG_DIR, "history.csv")
tb_logdir = os.path.join(LOG_DIR, "tensorboard")

callbacks = [
    # Full model (architecture + weights + optimizer state)
    tf.keras.callbacks.ModelCheckpoint(best_full, monitor="val_accuracy", mode="max",
                                       save_best_only=True, save_weights_only=False),
    tf.keras.callbacks.ModelCheckpoint(last_full, monitor="val_accuracy", mode="max",
                                       save_best_only=False, save_weights_only=False),

    # Weights-only (if someone wants to rebuild model code)
    tf.keras.callbacks.ModelCheckpoint(best_wts, monitor="val_accuracy", mode="max",
                                       save_best_only=True, save_weights_only=True),
    tf.keras.callbacks.ModelCheckpoint(last_wts, monitor="val_accuracy", mode="max",
                                       save_best_only=False, save_weights_only=True),

    tf.keras.callbacks.CSVLogger(csv_log, append=True),
    tf.keras.callbacks.TensorBoard(log_dir=tb_logdir),
    tf.keras.callbacks.EarlyStopping(monitor="val_accuracy", mode="max", patience=8, restore_best_weights=True)
]

def get_initial_epoch_from_csv(csv_path: str) -> int:
    # If file doesn't exist or is empty -> start from epoch 0
    if (not os.path.exists(csv_path)) or os.path.getsize(csv_path) == 0:
        return 0

    try:
        df = pd.read_csv(csv_path)
        # If somehow it has headers but no rows, still 0
        return int(len(df))
    except pd.errors.EmptyDataError:
        return 0
    except Exception as e:
        # If it's corrupted, safest is to restart or rename it
        print(f"[WARN] Could not read CSV log ({csv_path}): {e}. Starting from epoch 0.")
        return 0

def maybe_resume_or_build(input_shape, num_classes, lr):
    initial_epoch = get_initial_epoch_from_csv(csv_log)

    if os.path.exists(last_full):

```

```

print("Resuming from:", last_full)
model = tf.keras.models.load_model(last_full)
print("Initial epoch:", initial_epoch)
return model, initial_epoch

print("No checkpoint found. Building new model.")
model = build_cnn(input_shape, num_classes, lr)
return model, 0

```

PART 1: CIFAR-10 (Keras)

```

if PART == 1:
    (x_train, y_train), (x_test, y_test) = tf.keras.datasets.cifar10.load_data()
    y_train = y_train.squeeze()
    y_test = y_test.squeeze()

    # Keep as uint8 for Rescaling(1./255) layer
    x_train = x_train.astype("uint8")
    x_test = x_test.astype("uint8")

    print("Train:", x_train.shape, y_train.shape)
    print("Test :", x_test.shape, y_test.shape)

    # Visualize some images:
    show_samples_np(x_train, y_train, n=12)

    input_shape = (32, 32, 3)
    num_classes = 10

    model, initial_epoch = maybe_resume_or_build(input_shape, num_classes, CONFIG["lr"])

    # Save model summary (useful for grading)
    with open(os.path.join(RUN_DIR, "model_summary.txt"), "w") as f:
        model.summary(print_fn=lambda s: f.write(s + "\n"))

    history = model.fit(
        x_train, y_train,
        validation_data=(x_test, y_test),
        epochs=CONFIG["epochs"],
        initial_epoch=initial_epoch,
        batch_size=CONFIG["batch_size"],
        callbacks=callbacks,
        verbose=1
    )

    # Visualize model accuracy over training epochs for both training and validation datasets
    plot_accuracy_from_csv(csv_log, title="Part 1 (CIFAR-10) Accuracy vs Epoch")


    # -----
    # Display model predictions on test images
    # Shows predicted label (p) vs true label (t)
    # Used to visually evaluate model performance
    # -----
    preds = model.predict(x_test[:32], verbose=0).argmax(axis=1)
    plt.figure(figsize=(12,4))
    for i in range(12):
        plt.subplot(2, 6, i+1)
        plt.imshow(x_test[i])
        plt.title(f"p:{preds[i]} / t:{y_test[i]}")
        plt.axis("off")
    plt.tight_layout()
    plt.show()

```

Train: (50000, 32, 32, 3) (50000,)
 Test : (10000, 32, 32, 3) (10000,)



Resuming from: /content/drive/MyDrive/DL_Prog2/part1_run1/checkpoints/last.keras
 Initial epoch: 30



Print Best Training and Validation Accuracy

```
# -----
# Print best accuracy values
# -----
df = pd.read_csv(csv_log)

if "accuracy" in df.columns:
    best_train_acc = df["accuracy"].max()
    final_train_acc = df["accuracy"].iloc[-1]
    print(f"Best Train Accuracy : {best_train_acc:.4f} ({best_train_acc*100:.2f}%)")
```

```
print(f"Final Train Accuracy : {final_train_acc:.4f} `({final_train_acc*100:.2f}%)"

if "val_accuracy" in df.columns:
    best_val_acc = df["val_accuracy"].max()
    final_val_acc = df["val_accuracy"].iloc[-1]
    print(f"Best Validation Accuracy : {best_val_acc:.4f} ({best_val_acc*100:.2f}%)"
    print(f"Final Validation Accuracy : {final_val_acc:.4f} ({final_val_acc*100:.2f}%)"

Best Train Accuracy : 0.7649 (76.49%)
Final Train Accuracy : 0.7649 (76.49%)
Best Validation Accuracy : 0.7857 (78.57%)
Final Validation Accuracy : 0.7840 (78.40%)
```

PART 2: CIFAR-100 (Keras)

```
# Select which assignment part to run
PART = 2      # 1 = CIFAR10
              # 2 = CIFAR100
              # 3 = Animal10

RUN_NAME = f"part{PART}_run1"
```

```
from google.colab import drive
drive.mount('/content/drive')

import os

PROJECT_DIR = "/content/drive/MyDrive/DL_Prog2"
RUN_NAME = "part2_run1"

RUN_DIR = os.path.join(PROJECT_DIR, RUN_NAME)

CHECKPOINT_DIR = os.path.join(RUN_DIR, "checkpoints")
LOG_DIR = os.path.join(RUN_DIR, "logs")

os.makedirs(CHECKPOINT_DIR, exist_ok=True)
os.makedirs(LOG_DIR, exist_ok=True)

print("Checkpoint directory:", CHECKPOINT_DIR)
print("Log directory:", LOG_DIR)
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True)

Checkpoint directory: /content/drive/MyDrive/DL_Prog2/part2_run1/checkpoints

Log directory: /content/drive/MyDrive/DL_Prog2/part2_run1/logs

```
if PART == 2:
    (x_train, y_train), (x_test, y_test) = tf.keras.datasets.cifar100.load_data()
    y_train = y_train.squeeze()
    y_test = y_test.squeeze()

    # Keep as uint8 for Rescaling(1./255) layer
    x_train = x_train.astype("uint8")
    x_test = x_test.astype("uint8")

    print("Train:", x_train.shape, y_train.shape)
    print("Test :", x_test.shape, y_test.shape)

    # Visualize some images (required)
    show_samples_np(x_train, y_train, n=12)

    input_shape = (32, 32, 3)
    num_classes = 100

    model, initial_epoch = maybe_resume_or_build(input_shape, num_classes, CONFIG["lr"])

    # Save model summary (useful for grading)
    with open(os.path.join(RUN_DIR, "model_summary.txt"), "w") as f:
        model.summary(print_fn=lambda s: f.write(s + "\n"))

    history = model.fit(
        x_train, y_train,
        validation_data=(x_test, y_test),
        epochs=CONFIG["epochs"],
        initial_epoch=initial_epoch,
```

```
        batch_size=CONFIG["batch_size"],
        callbacks=callbacks,
        verbose=1
    )



    # Plot accuracy over epochs (required)
    plot_accuracy_from_csv(csv_log, title="Part 2 (CIFAR-100) Accuracy vs Epoch")

    # Show some predictions (required)
    preds = model.predict(x_test[:32], verbose=0).argmax(axis=1)
    plt.figure(figsize=(12,4))
    for i in range(12):
        plt.subplot(2, 6, i+1)
        plt.imshow(x_test[i])
        plt.title(f"p:{preds[i]} / t:{y_test[i]}")
        plt.axis("off")
    plt.tight_layout()
    plt.show()
```

Train: (50000, 32, 32, 3) (50000,)
 Test : (10000, 32, 32, 3) (10000,)



Resuming from: /content/drive/MyDrive/DL_Prog2/part1_run1/checkpoints/last.keras
 Initial epoch: 30



Print Best Training and Validation Accuracy

```
# -----
# Print best accuracy values
# -----
df = pd.read_csv(csv_log)

if "accuracy" in df.columns:
    best_train_acc = df["accuracy"].max()
    final_train_acc = df["accuracy"].iloc[-1]
```



```

print(f"Best Train Accuracy : {best_train_acc:.4f} ({best_train_acc*100:.2f}%)")
print(f"Final Train Accuracy : {final_train_acc:.4f} ({final_train_acc*100:.2f}%)")

if "val_accuracy" in df.columns:
    best_val_acc = df["val_accuracy"].max()
    final_val_acc = df["val_accuracy"].iloc[-1]
    print(f"Best Validation Accuracy : {best_val_acc:.4f} ({best_val_acc*100:.2f}%)")
    print(f"Final Validation Accuracy : {final_val_acc:.4f} ({final_val_acc*100:.2f}%)")

Best Train Accuracy : 0.7649 (76.49%)
Final Train Accuracy : 0.7649 (76.49%)
Best Validation Accuracy : 0.7857 (78.57%)
Final Validation Accuracy : 0.7840 (78.40%)

```

PART 3: Animal10

```

# Select which assignment part to run
PART = 3      # 1 = CIFAR10
              # 2 = CIFAR100
              # 3 = Animal10

```

```
RUN_NAME = f"part{PART}_run1"
```

```

from google.colab import drive
drive.mount('/content/drive')

import os

# controller
# PART = 3
RUN_NAME = f"part{PART}_run1"

PROJECT_DIR = "/content/drive/MyDrive/DL_Prog2"
RUN_DIR = os.path.join(PROJECT_DIR, RUN_NAME)

CKPT_DIR = os.path.join(RUN_DIR, "checkpoints")
LOG_DIR = os.path.join(RUN_DIR, "logs")
OUT_DIR = os.path.join(RUN_DIR, "outputs")

for d in [CKPT_DIR, LOG_DIR, OUT_DIR]:
    os.makedirs(d, exist_ok=True)

print("RUN_DIR:", RUN_DIR)
print("CKPT_DIR:", CKPT_DIR)
print("LOG_DIR :", LOG_DIR)
print("OUT_DIR :", OUT_DIR)

```

```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True)
RUN_DIR: /content/drive/MyDrive/DL_Prog2/part3_run1
CKPT_DIR: /content/drive/MyDrive/DL_Prog2/part3_run1/checkpoints
LOG_DIR : /content/drive/MyDrive/DL_Prog2/part3_run1/logs
OUT_DIR : /content/drive/MyDrive/DL_Prog2/part3_run1/outputs

```

```

if PART == 3:
    import os, json
    import numpy as np
    import pandas as pd
    import tensorflow as tf
    import matplotlib.pyplot as plt

    # -----
    # 1) Dataset path (yours)
    # -----
    ANIMAL_DIR = "/content/drive/MyDrive/AnimalsData/Animals10/raw-img"
    assert os.path.exists(ANIMAL_DIR), f"Not found: {ANIMAL_DIR}"

    IMG_SIZE = CONFIG.get("img_size_part3", 224)
    BATCH = CONFIG["batch_size"]
    VAL_SPLIT = CONFIG.get("val_split", 0.2)
    SEED = CONFIG.get("seed", 42)

    AUTOTUNE = tf.data.AUTOTUNE

    # -----
    # 2) Load with image_dataset_from_directory

```

```

# -----
train_ds = tf.keras.utils.image_dataset_from_directory(
    ANIMAL_DIR,
    validation_split=VAL_SPLIT,
    subset="training",
    seed=SEED,
    image_size=(IMG_SIZE, IMG_SIZE),
    batch_size=BATCH,
    label_mode="int"
)
val_ds = tf.keras.utils.image_dataset_from_directory(
    ANIMAL_DIR,
    validation_split=VAL_SPLIT,
    subset="validation",
    seed=SEED,
    image_size=(IMG_SIZE, IMG_SIZE),
    batch_size=BATCH,
    label_mode="int"
)

class_names = train_ds.class_names
num_classes = len(class_names)
print("Classes:", class_names)
print("num_classes:", num_classes)

# Save class mapping for grading
with open(os.path.join(RUN_DIR, "class_names.json"), "w") as f:
    json.dump(class_names, f, indent=2)

# -----
# 3) Explore: show a batch + distribution
# -----
def show_ds_samples(ds, class_names, n=12):
    plt.figure(figsize=(12,4))
    for images, labels in ds.take(1):
        print("Batch:", images.shape, images.dtype)
        for i in range(min(n, images.shape[0])):
            plt.subplot(2, n//2, i+1)
            plt.imshow(images[i].numpy().astype("uint8"))
            plt.title(class_names[int(labels[i])])
            plt.axis("off")
    plt.tight_layout()
    plt.show()

def class_distribution(ds, num_classes):
    counts = np.zeros(num_classes, dtype=np.int64)
    for _, labels in ds:
        labels_np = labels.numpy()
        for c in labels_np:
            counts[int(c)] += 1
    return counts

show_ds_samples(train_ds, class_names, n=12)

counts = class_distribution(train_ds, num_classes)
print("Train counts:")
for name, c in zip(class_names, counts.tolist()):
    print(f"{name:12s}: {c}")

# Class weights (helps imbalance without complicated sampling)
total = counts.sum()
class_weight = {i: (total / (num_classes * max(counts[i], 1))) for i in range(num_classes)}
print("Class weights:", class_weight)

# Speed pipeline
train_ds = train_ds.cache().shuffle(2000, seed=SEED).prefetch(AUTOTUNE)
val_ds = val_ds.cache().prefetch(AUTOTUNE)

# -----
# 4) Transfer learning model (MobileNetV2)
# -----
# Data augmentation
data_augmentation = tf.keras.Sequential([
    tf.keras.layers.RandomFlip("horizontal"),
    tf.keras.layers.RandomRotation(0.08),
    tf.keras.layers.RandomZoom(0.15),
    tf.keras.layers.RandomContrast(0.15),
    tf.keras.layers.RandomTranslation(0.05, 0.05),
])

def build_transfer_model(num_classes, img_size, lr, base_trainable=False):
    inputs = tf.keras.Input(shape=(img_size, img_size, 3))

```

```

x = data_augmentation(inputs)

# MobileNetV2 expects pixel values in [-1, 1]
x = tf.keras.applications.mobilenet_v2.preprocess_input(x)

base = tf.keras.applications.MobileNetV2(
    include_top=False,
    weights="imagenet",
    input_shape=(img_size, img_size, 3)
)
base.trainable = base_trainable # False for feature extraction phase

x = base(x, training=False)
x = tf.keras.layers.GlobalAveragePooling2D()(x)
x = tf.keras.layers.Dropout(0.3)(x)
x = tf.keras.layers.Dense(256, activation="relu")(x)
x = tf.keras.layers.Dropout(0.2)(x)
outputs = tf.keras.layers.Dense(num_classes, activation="softmax")(x)

model = tf.keras.Model(inputs, outputs)
model.compile(
    optimizer=tf.keras.optimizers.Adam(lr),
    loss=tf.keras.losses.SparseCategoricalCrossentropy(),
    metrics=["accuracy"]
)
return model

# -----
# 5) Resume helpers (Part 3)
# -----
last_full = os.path.join(CKPT_DIR, "last.keras")
best_full = os.path.join(CKPT_DIR, "best.keras")
csv_log = os.path.join(LOG_DIR, "history.csv")

def initial_epoch_from_csv(path):
    if not os.path.exists(path):
        return 0
    df = pd.read_csv(path)
    return len(df)

# If a checkpoint exists, load it; otherwise build new
if os.path.exists(last_full):
    print("Resuming from:", last_full)
    model = tf.keras.models.load_model(last_full)
    initial_epoch = initial_epoch_from_csv(csv_log)
    print("Initial epoch:", initial_epoch)
else:
    # Phase 1: feature extraction (base frozen)
    model = build_transfer_model(num_classes, IMG_SIZE, lr=CONFIG["lr"], base_trainable=False)
    initial_epoch = 0

# Save model summary
with open(os.path.join(RUN_DIR, "model_summary.txt"), "w") as f:
    model.summary(print_fn=lambda s: f.write(s + "\n"))

# -----
# 6) Callbacks (best/last + logs) - reuse yours if already defined
# -----
callbacks = [
    tf.keras.callbacks.ModelCheckpoint(best_full, monitor="val_accuracy", mode="max",
                                       save_best_only=True, save_weights_only=False),
    tf.keras.callbacks.ModelCheckpoint(last_full, monitor="val_accuracy", mode="max",
                                       save_best_only=False, save_weights_only=False),
    tf.keras.callbacks.CSVLogger(csv_log, append=True),
    tf.keras.callbacks.ReduceLROnPlateau(monitor="val_loss", factor=0.3, patience=2, min_lr=1e-7, verbose=1),
    tf.keras.callbacks.EarlyStopping(monitor="val_accuracy", mode="max", patience=6,
                                     restore_best_weights=True),
]

# -----
# 7) Train Phase 1 (frozen base)
# -----
EPOCHS_TOTAL = CONFIG["epochs"]
phase1_epochs = min(10, EPOCHS_TOTAL) # quick head training

print("\n=== Phase 1: Frozen base ===")
history1 = model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=phase1_epochs,
    initial_epoch=initial_epoch,
    callbacks=callbacks,

```

```

class_weight=class_weight,
verbose=1
)

# -----
# 8) Fine-tuning Phase 2 (unfreeze last layers)
# -----
# Only fine-tune if we still have epochs left
current_epoch = initial_epoch_from_csv(csv_log)
if current_epoch < EPOCHS_TOTAL:
    print("\n=== Phase 2: Fine-tuning ===")

    # Unfreeze the base model if present
    # Find MobileNetV2 layer by type/name
    base_model = None
    for layer in model.layers:
        if isinstance(layer, tf.keras.Model) and "mobilenetv2" in layer.name.lower():
            base_model = layer
            break

    if base_model is not None:
        base_model.trainable = True

        # Fine-tune only the last N layers (keeps training stable)
        fine_tune_at = max(0, len(base_model.layers) - 60)
        for l in base_model.layers[:fine_tune_at]:
            l.trainable = False

        # Use smaller LR for fine-tuning
        model.compile(
            optimizer=tf.keras.optimizers.Adam(learning_rate=1e-5),
            loss=tf.keras.losses.SparseCategoricalCrossentropy( ),
            metrics=["accuracy"]
        )

    history2 = model.fit(
        train_ds,
        validation_data=val_ds,
        epochs=EPOCHS_TOTAL,
        initial_epoch=current_epoch,
        callbacks=callbacks,
        class_weight=class_weight,
        verbose=1
    )

# -----
# 9) Plot accuracy vs epochs
# -----
df = pd.read_csv(csv_log)
epochs_axis = np.arange(len(df))

plt.figure()
if "accuracy" in df.columns:
    plt.plot(epochs_axis, df["accuracy"], label="train_acc")
if "val_accuracy" in df.columns:
    plt.plot(epochs_axis, df["val_accuracy"], label="val_acc")
plt.xlabel("Epoch")
plt.ylabel("Accuracy")
plt.title("Part 3 (Animal10) Transfer Learning Accuracy vs Epoch")
plt.legend()
plt.grid(True)
plt.show()

# -----
# 10) Show some predictions
# -----
best_model = tf.keras.models.load_model(best_full)
for images, labels in val_ds.take(1):
    preds = best_model.predict(images, verbose=0).argmax(axis=1)
    plt.figure(figsize=(12,4))
    for i in range(12):
        plt.subplot(2, 6, i+1)
        plt.imshow(images[i].numpy().astype("uint8"))
        plt.title(f"{class_names[int(preds[i])]} / {class_names[int(labels[i])]}")
        plt.axis("off")
    plt.tight_layout()
    plt.show()

print("\nSaved:")
print(" - best:", best_full)
print(" - last:", last_full)

```

```
print(" - log :", csv_log)
print(" - class_names.json + config.json + model_summary.txt in RUN_DIR")
```